

Virtual Memory


- Background
- Demand Paging
- Process Creation
- Page Replacement
- Allocation of Frames
- Thrashing
- Operating System Examples

Background

- **Virtual memory** – separation of user logical memory from physical memory.
 - ☞ Only part of the program needs to be in memory for execution.
 - ☞ Logical address space can therefore be much larger than physical address space.
 - ☞ Allows address spaces to be shared by several processes.
 - ☞ Allows for more efficient process creation.

- Virtual memory can be implemented via:
 - ☞ Demand paging
 - ☞ Demand segmentation


Virtual Memory That is Larger Than Physical Memory



Demand Paging

- Bring a page into memory only when it is needed.
 - ☞ Less I/O needed
 - ☞ Less memory needed
 - ☞ Faster response
 - ☞ More users
- Page is needed \Rightarrow reference to it
 - ☞ invalid reference \Rightarrow abort
 - ☞ not-in-memory \Rightarrow bring to memory

Transfer of a Paged Memory to Contiguous Disk Space



Valid-Invalid Bit

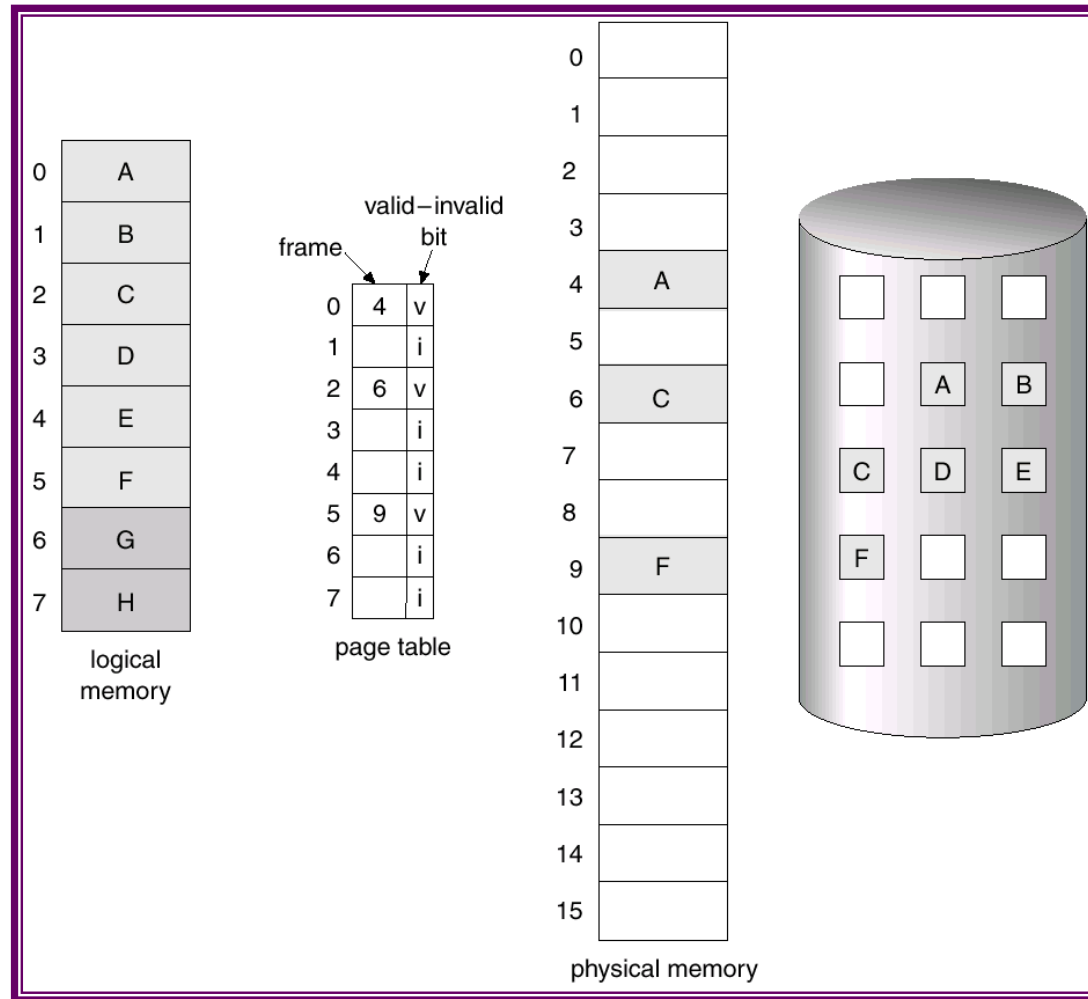
- With each page table entry a valid–invalid bit is associated
(1 \Rightarrow in-memory, 0 \Rightarrow not-in-memory)
- Initially valid–invalid bit is set to 0 on all entries.
- Example of a page table snapshot.

Frame #	valid-invalid bit
	1
	1
	1
	1
	0
⋮	
	0
	0

page table


- During address translation, if valid–invalid bit in page table entry is 0 \Rightarrow page fault.

Page Table When Some Pages Are Not in Main Memory



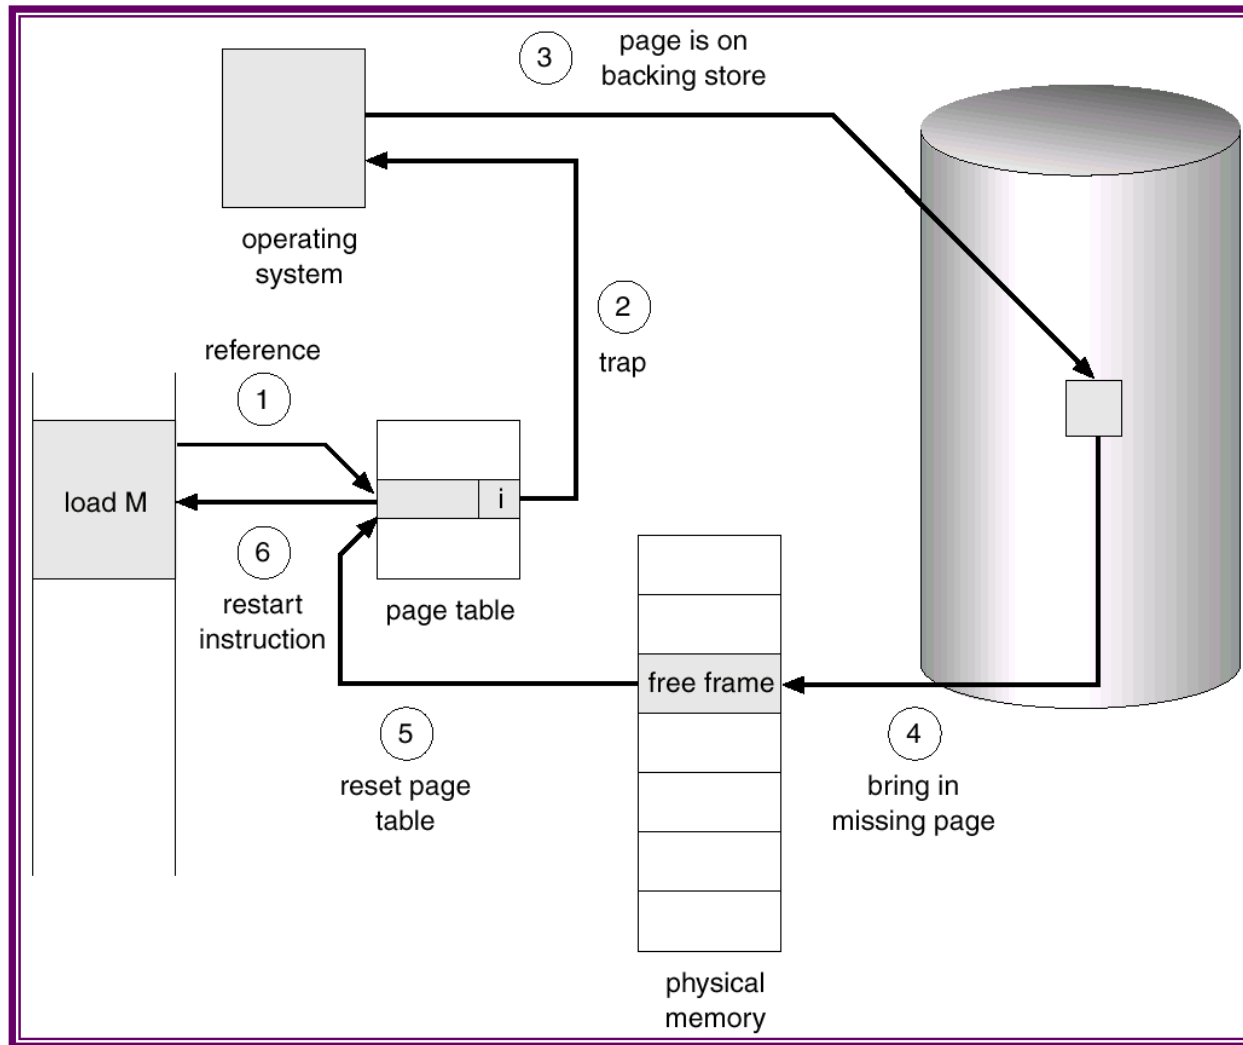
Page Fault

- If there is ever a reference to a page, first reference will trap to OS \Rightarrow page fault
- OS looks at another table to decide:
 - ☞ Invalid reference \Rightarrow abort.
 - ☞ Just not in memory.
- Get empty frame.
- Swap page into frame.
- Reset tables, validation bit = 1.
- Restart instruction: Least Recently Used
 - ☞ block move



☞ auto increment/decrement location

Steps in Handling a Page Fault





Performance of Demand Paging

- Stages in Demand Paging
 1. Trap to the operating system
 2. Save the user registers and process state
 3. Determine that the interrupt was a page fault
 4. Check that the page reference was legal and determine the location of the page on the disk
 5. Issue a read from the disk to a free frame:
 1. Wait in a queue for this device until the read request is serviced
 2. Wait for the device seek and/or latency time
 3. Begin the transfer of the page to a free frame
 6. While waiting, allocate the CPU to some other user
 7. Receive an interrupt from the disk I/O subsystem (I/O completed)
 8. Save the registers and process state for the other user
 9. Determine that the interrupt was from the disk
 10. Correct the page table and other tables to show page is now in memory
 11. Wait for the CPU to be allocated to this process again
 12. Restore the user registers, process state, and new page table, and then resume the interrupted instruction



Performance of Demand Paging

■ Page Fault Rate $0 \leq p \leq 1.0$

☞ if $p = 0$ no page faults

☞ if $p = 1$, every reference is a fault

■ Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & + [\text{swap page out}] \\ & + \text{swap page in} \\ & + \text{restart overhead}) \end{aligned}$$


Demand Paging Example

- Memory access time = 200 nanosecond
- 50% of the time the page that is being replaced has been modified and therefore needs to be swapped out.
- Page fault Time = 8 msec = 8,000,000
$$EAT = (1 - p) \times 200 + p (8,000,000)$$

What happens if there is no free frame?

- Page replacement – find some page in memory, but not really in use, swap it out.
 - ☞ algorithm
 - ☞ performance – want an algorithm which will result in minimum number of page faults.
- Same page may be brought into memory several times.

Need For Page Replacement



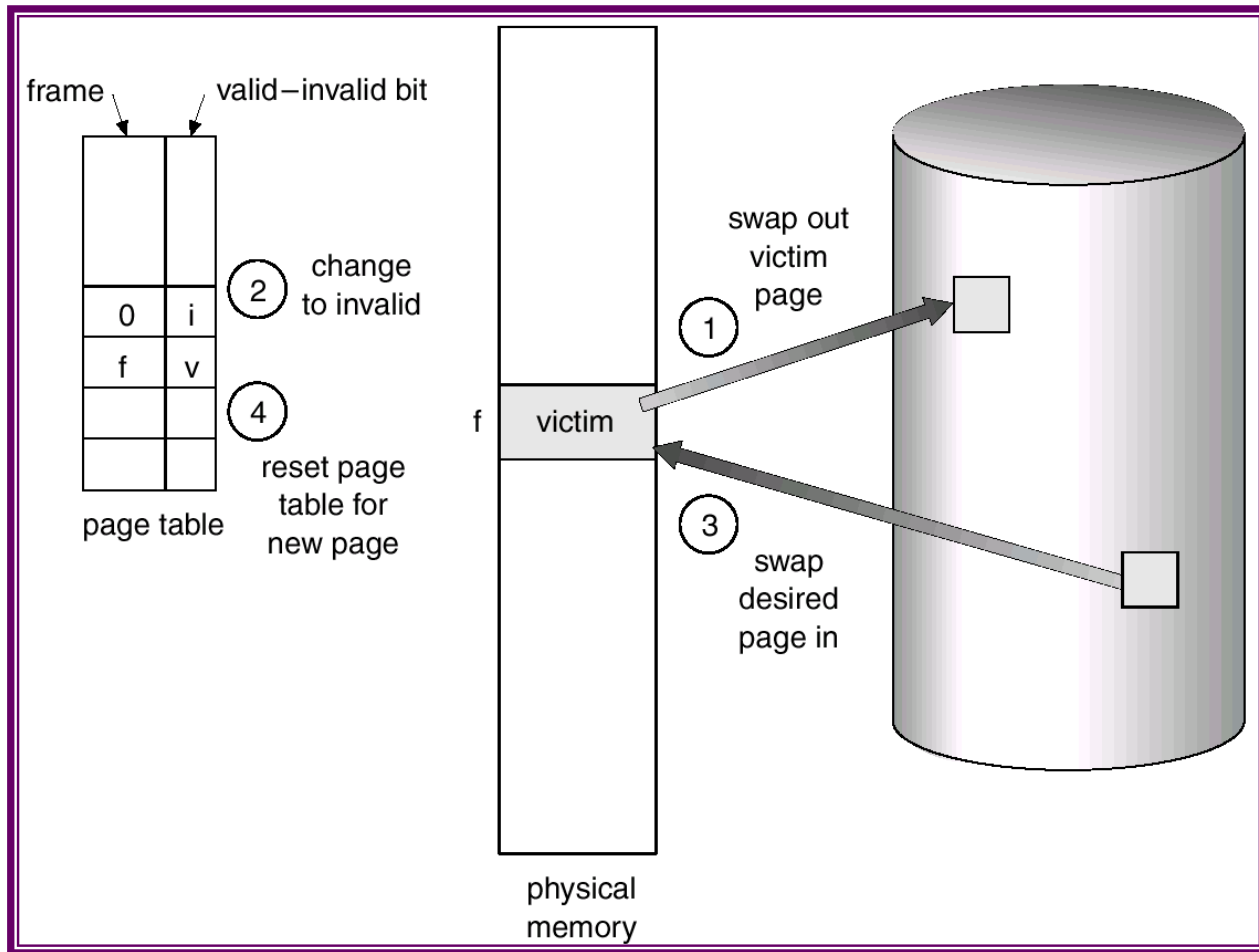
Page Replacement

- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement.
- Use *modify (dirty) bit* to reduce overhead of page transfers – only modified pages are written to disk.
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory.

Basic Page Replacement

1. Find the location of the desired page on disk.
2. Find a free frame:
 - If there is a free frame, use it.
 - If there is no free frame, use a page replacement algorithm to select a *victim* frame.
3. Read the desired page into the (newly) free frame.
Update the page and frame tables.
4. Restart the process.


Page Replacement




Page Replacement Algorithms

- Want lowest page-fault rate.
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string.
- In all our examples, the reference string is
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5.

Graph of Page Faults Versus The Number of Frames



FIFO Page Replacement



First-In-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)


1	1	4	5	
2	2	1	3	9 page faults
3	3	2	4	

- 4 frames

1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		

- FIFO Replacement – Belady's Anomaly
☞ more frames \Rightarrow less page faults


FIFO Illustrating Belady's Anomaly



Optimal Algorithm

- Replace page that will not be used for longest period of time.
- 4 frames example


1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



6 page faults

- How do you know this?
- Used for measuring how well your algorithm performs.

Optimal Page Replacement



Least Recently Used (LRU) Algorithm


- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	5	
2		
3	5	4
4	3	

- Counter implementation

- ☞ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter.
- ☞ When a page needs to be changed, look at the counters to determine which are to change.


LRU Page Replacement



LRU Algorithm (Cont.)

- Stack implementation – keep a stack of page numbers in a double link form:
 - ☞ Page referenced:
 - 📄 move it to the top
 - 📄 requires 6 pointers to be changed
 - ☞ No search for replacement

Use Of A Stack To Record The Most Recent Page References



LRU Approximation Algorithms

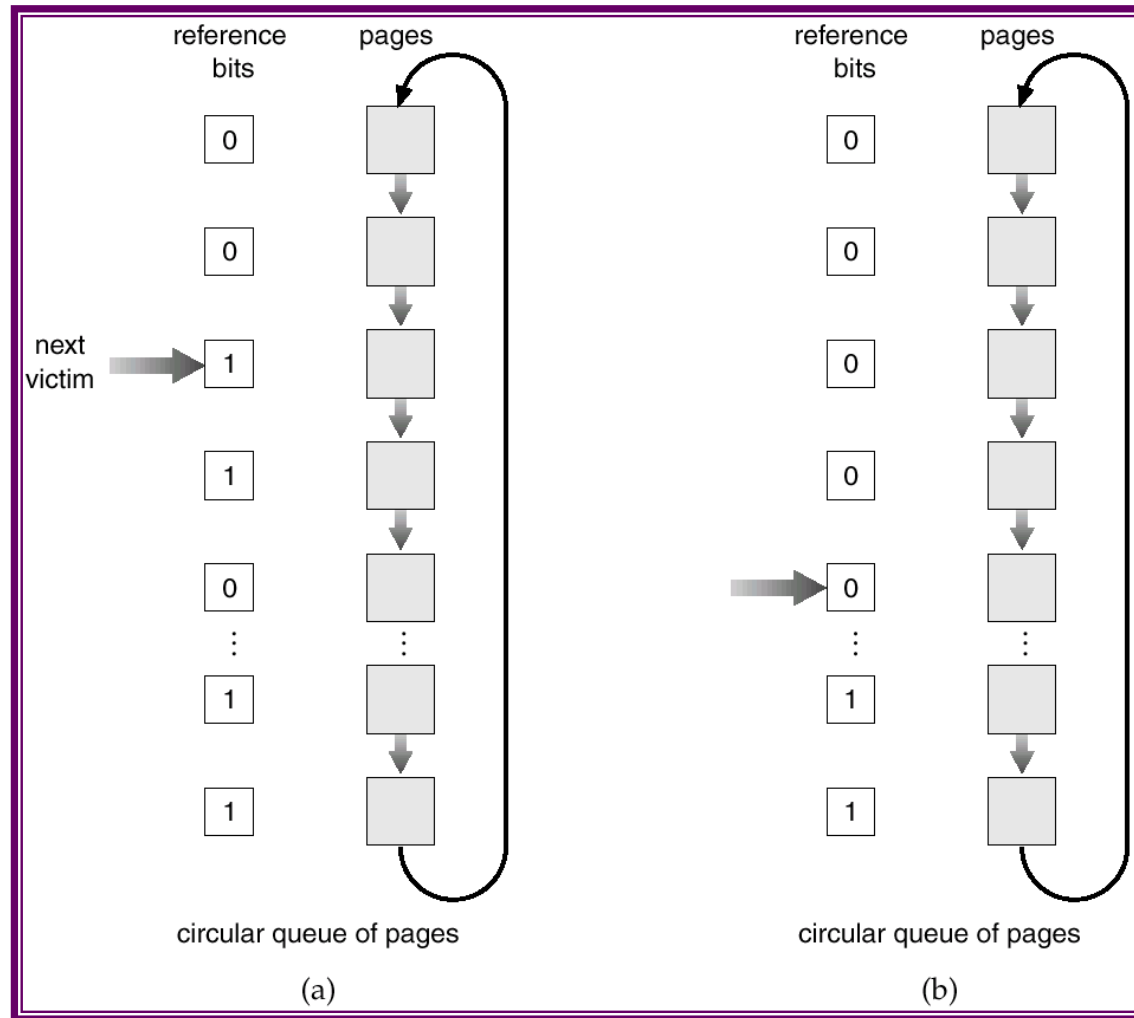
■ Reference bit

- ☞ With each page associate a bit, initially = 0
- ☞ When page is referenced bit set to 1.
- ☞ Replace the one which is 0 (if one exists). We do not know the order, however.

■ Second chance

- ☞ Need reference bit.
- ☞ Clock replacement.
- ☞ If page to be replaced (in clock order) has reference bit = 1. then:
 - 📄 set reference bit 0.
 - 📄 leave page in memory.
 - 📄 replace next page (in clock order), subject to same rules.

Second-Chance (clock) Page-Replacement Algorithm



Counting Algorithms

- Keep a counter of the number of references that have been made to each page.
- LFU Algorithm: replaces page with smallest count.
- MFU Algorithm: based on the argument that the page with the smallest count was probably just brought in and has yet to be used.

Allocation of Frames

- Each process needs **minimum** number of pages.
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - ☞ instruction is 6 bytes, might span 2 pages.
 - ☞ 2 pages to handle **from**.
 - ☞ 2 pages to handle **to**.
- Two major allocation schemes.
 - ☞ fixed allocation
 - ☞ priority allocation

Fixed Allocation

- Equal allocation – e.g., if 100 frames and 5 processes, give each 20 pages.
- Proportional allocation – Allocate according to the size of process.

– s_i = size of process p_i

– $S = \sum s_i$

– m = total number of frames

– a_i = allocation for $p_i = \frac{s_i}{S} \times m$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

Priority Allocation

- Use a proportional allocation scheme using priorities rather than size.
- If process P_i generates a page fault,
 - ☞ select for replacement one of its frames.
 - ☞ select for replacement a frame from a process with lower priority number.


Global vs. Local Allocation

- **Global** replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another.
- **Local** replacement – each process selects from only its own set of allocated frames.

Thrashing


- If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
 - ☞ low CPU utilization.
 - ☞ operating system thinks that it needs to increase the degree of multiprogramming.
 - ☞ another process added to the system.
- **Thrashing** \equiv a process is busy swapping pages in and out.

Thrashing



- Why does paging work?
Locality model
 - ☞ Process migrates from one locality to another.
 - ☞ Localities may overlap.
- Why does thrashing occur?
 Σ size of locality > total memory size


Locality In A Memory-Reference Pattern




Working-Set Model

- $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instruction
- WSS_i (working set of Process P_i) =
total number of pages referenced in the most recent Δ
(varies in time)
 - ☞ if Δ too small will not encompass entire locality.
 - ☞ if Δ too large will encompass several localities.
 - ☞ if $\Delta = \infty \Rightarrow$ will encompass entire program.
- $D = \sum WSS_i \equiv$ total demand frames
- if $D > m \Rightarrow$ Thrashing
- Policy if $D > m$, then suspend one of the processes.

Working-set model



Page-Fault Frequency Scheme



- Establish “acceptable” page-fault rate.
 - ☞ If actual rate too low, process loses frame.
 - ☞ If actual rate too high, process gains frame.